



AI Workbench

it
works

How

Grading

Compare

FAQ

Docs

Run it locally

> AI Workbench

MCP ANALYZER AND SESSION WORKBENCH

Trust every AI session before it ships.

AI Workbench measures what your MCP servers and Agent Skills cost in context, drives them through real sessions, grades every run automatically, and compares anything. That includes a Qlik Answers assistant against a full MCP agent. It runs on your own machine.

Run it locally

See how it works

RUN, GRADED

Is American Airlines still our best partner, or should we switch?

Answer: adopt a dual-carrier strategy

TURNS	COST	QUALITY
9	\$0.31	100%



Answer validated

Cited evidence

0 errors

WHY IT MATTERS

Your context budget is spent before the model says a word.

Every tool and skill definition loads into the model's context up front. Sessions then fail in ways you cannot see, and you pick between agents and assistants on instinct. AI Workbench turns all of it into numbers.

231,499

startup tokens
measured across one
team's seven connected
servers, before a single
question is asked.

Measured live in AI
Workbench.

9 to 1

turns to answer the
same question: a full
MCP agent versus a Qlik
Answers assistant. One
costs \$0.31. The other
costs nothing.

Same test, two run targets,
compared side by side.

100% vs 70%

the quality trade-off
between those two
approaches, graded
automatically rather
than guessed at.

Answer-validation score per
run.

WHO IT HELPS

One setup, read four different ways.

The same runs answer a different question for everyone who touches the setup. Nobody has to guess from a log again.

OPERATORS

MCP OWNERS

Understand your setup

See what your servers and skills actually cost in context, and where a session's budget goes before the model answers.

PRESALES & CSES

Debug a real session

Replay any run turn by turn to show where it went wrong, why, and what a fix would change, with numbers to back it.

SKILL & MCP DEVELOPERS

Find issues and fix them

Automatic issue detection files a tracked bug with a drafted fix, so a failing run becomes a test-and-fix loop.

Validate and track servers

Analyze a server, watch its footprint change across versions, and run it as a long-running quality gate.

MEASURE THE FOOTPRINT

Start by seeing what your setup costs.

Connect a server and scan it. Every tool is ranked by the real tokens its definition loads into context, so you know what you are paying for before a single question is asked.

Startup footprint, 7 servers



SERVERS

TOKENS

RANKED

7

231K by cost



■ Heaviest server


■ Lightest

Ranked by real token cost


Each tool and skill definition is measured with production tokenizers, then sorted heaviest first with findings on what to trim.

Every server in one view

Connected servers sit side by side with their tool counts and combined startup cost, updated as you add or remove them.



Footprint scan. Every tool ranked by the real tokens it loads into context, heaviest first.



Connected servers. Seven servers measured together: 231,499 startup tokens before any question.

The whole session, not just the answer.

Every run is captured start to finish. Read it like a transcript, or drop into an event-level trace when something goes wrong.

Live KPI rail

Context used against the model's limit, estimated cost, provider-actual tokens, tool calls, and turns, visible on every run.

Context, turn by turn

See exactly what fills the window each turn: system, tool definitions, history, tool results, and output.

Trace to the exact step


Expand any turn into its prompts, responses, and tool calls, each stamped with tokens, duration, and cost.


Context window, 9 turns




CONTEXT COST TOKENS

2% **\$0.31** **261K**





Run console. The whole session as a transcript, with context, cost, tokens, and turns on the KPI rail.




Event trace. Expand any turn into its prompts, responses, and tool calls, each stamped with tokens and time.

BRING YOUR OWN MODEL

Run the same session on any model you choose.


AI Workbench drives every provider through the same agent loop, so you can compare them like for like — in the cloud, on your own machine, or against Qlik's own assistant.



THE VALIDATION LOOP

Define once, test everywhere, feed the fixes back.

Comparing models by hand means re-running the same questions again and again and eyeballing every answer. The validation loop replaces that: set up a test once, run it across every model you support, compare the results in a single view, then send validated fixes — a tracked issue with a drafted change — straight back to the MCP server and skill that need them.



Models tested: Anthropic · Anthropic CLI · OpenAI · Google · Ollama · Qlik Answers

GRADE EVERY RUN

Every run is graded, automatically.

The moment a run finishes it is rated, with no reference answer required. That grade is what turns a pile of sessions into something you can rank, trust, and compare.


Run rating



Answered

Surplus: valuable

0 error findings



CITED EVIDENCE

"AA ranks 6th of 8 major carriers on arrival punctuality, and 5th of 8 on cancellations."

ROUGE-1 HYGIENE JUDGE ANSWER SURPLUS

8% 100% 80% 100% 90%

Answer validation, with evidence


Did the answer address the task? You get a verdict and a 0 to 1 score, backed by quotes lifted straight from the transcript. No black-box numbers.

Insight surplus and error forensics


Useful extra content lifts the score. Padding lowers it. A deterministic inventory catches every error, failed call, and guardrail stop.

Expectation grades when you want them

Add expectations and every run is also scored on Answer, Judge, ROUGE-1, Hygiene, and Surplus.



Run report. The automatic rating with its verdict, cited evidence, and every expectation grade in one place.



Run analytics. Where the tokens and cost went, turn by turn and category by category.

MEASURED, NOT GUESSED

Exactly what gets scored on every run.

Every run carries the same KPIs and the same grades, so any two runs sit on identical axes. That is what makes the comparisons trustworthy instead of anecdotal.

KPIs tracked on every run

Operational metrics, captured automatically as the session runs.

Turns and tool calls How much work the session did

Tokens in / out Provider-actual, not estimated

Cost Estimated from real model pricing

Peak context Tokens and % of the model's limit

Cached % Share served from prompt cache

Tool errors and duration Failures and wall-clock time

Questions Qlik Answers cost unit

How each run is graded

An automatic base rating, plus expectation grades when a test defines them.

Answer validation Verdict + 0 to 1 score, cited evidence

Insight surplus Value added vs. padding

Error forensics Root-cause inventory + drafted fix

ROUGE-1 Overlap with a reference answer

Judge Logprob-weighted LLM judge

Hygiene Clean, well-formed tool use

Trajectory and SkillFlow Path and skill conformance

Suite consistency Variance across repeated runs

Outcome

completed / aborted /
guardrail / failed

COMPARE ANYTHING

Qlik Answers against a full MCP agent, head to head.

Here is a comparison you cannot get anywhere else. AI Workbench treats a Qlik Answers assistant as a model, runs it against the same test as your MCP setup, and leads with a plain-language verdict.

Same question, apples to apples

Identical KPIs and grades on both sides mean the deltas

B vs A: 100% fewer tokens and 100% lower cost, at 30% lower quality. Dramatically cheaper and simpler, with a measurable trade-off.


RUN	TURNS	TOKENS	COST	QUALITY
A MCP agent claude-sonnet	9	266K	\$0.31	100%
B Qlik Answers assistant	1	478	\$0.00	70%
	89% less	100% less	100% less	30% less

Over time, model vs. model, skill on vs. off


The same workspace compares a run against its past, one model against another, or a skill's effect on a run.

See where sessions diverged

A turn-by-turn flow diff marks the exact step where two runs parted ways, down to the final answers side by side.



Side by side. Two runs on identical axes, with every KPI and grade delta called out.




Flow diff. A turn-by-turn view that marks the exact step where two sessions parted ways.


QLIK ANSWERS AS A MODEL

Run a Qlik Answers assistant like any other model.

Point a test at a Qlik Answers assistant and it runs as a first-class target: same agent loop, same KPIs, same grades. That is what makes the head-to-head against a full MCP agent an honest one, down to the cited evidence behind each answer.



A Qlik Answers run. One question, 478 tokens, graded on the same axes as a full MCP agent.



Head to head. The assistant against the MCP agent, with a plain-language verdict on the trade-off.

REAL ENGINE, REAL NUMBERS

Production tokenizers. A real agent loop. Zero cloud.

AI Workbench does not estimate and it does not proxy. It counts tokens the way the models do and drives servers through a real multi-provider agent loop, all inside one local container.

4

token profiles, including the real o200k and cl100k tokenizers, the same math the models use.

Every run, auto-graded

Answer validation, insight surplus, and error forensics run on every terminal session. No setup required.

100% local

Servers, credentials, and results stay on your machine. Secrets are encrypted at rest. No accounts, no cloud.

ONE WORKBENCH

Analyze, test, compare, and fix in one place.

Built for everyone who touches the setup: operators, Presales and CSEs, skill and MCP developers, and the owners who keep servers healthy.



Analyze

Scan any MCP server, rank every tool by real token cost, and get findings on where to recover context.



Test

Drive servers and skills through a real agent loop, with token accounting, cost estimates, and guardrails.



Compare

Put two sessions side by side: over time, model vs. model, skill on vs. off, or Qlik Answers vs. MCP.



Fix

Failing runs file tracked issues against your skills, with a drafted fix you can apply or hand to the assistant.

Qlik AI Workbench
MCP analyzer

Dashboard

MCP Servers

Scans

Compare

Skills

Skills

Testing

Collections

Runs

Compatibility

Setup

Environments

Settings

Local / dev mode

Since your last visit

Last visit Jul 13

No changes since Jul 13.

Needs attention

3

qlik-stage

Unauthorized

Failed

Scan now

Open

mcp-powerbi-fabric

No scan yet

Pending

Scan now

Open

mcp-assets

Connection refused at 192.168.65.254:7040 — nothing is accepting c...

Failed

Scan now

Open

Biggest movers

Largest token change since Jul 13

No footprint changes since Jul 13.

Servers

7

Total startup tokens

231,499

Tools + resources + prompts, latest successful scans

Resources

47

1,618 tokens

Prompts

0

0 tokens

Unscanned

1

Servers with no scan yet

Failed

2

Servers whose latest scan failed

Largest single tool

5,633

qlik_create_data_object

Tools scanned

272

Across latest successful scans

Latest server footprint

Server ↓	Tools ↓	Tokens ▼	Δ vs previous ↓	Largest tool ↓	Last scan ↓
qlik-stage	138	131,549		qlik_create_data_object	Jul 8, 2026, 2:58 PM
qlik-mreimtz	60	48,612	+3,348	qlik_create_data_object	Jul 12, 2026, 6:38 PM
barc-benchmark	60	48,612		qlik_create_data_object	Jul 12, 2026, 2:31 PM
mcp-assets	8	628		assets_update_metadata	Jun 22, 2026, 12:52 AM
mcp-textops	4	296		text_regex_extract	Jun 22, 2026, 12:05 AM
mcp-sql-local	2	184	+120	execute_sql	Jun 24, 2026, 12:57 PM

One dashboard. Servers, footprints, runs, and grades in a single local workspace.

THE LOOP CLOSES

Fixes flow back to the skill that caused them.

When a run grades poorly and a skill was involved, AI Workbench files a tracked issue on that skill, with a drafted fix and a link to the failing run. Resolve it, or hand it to the assistant for a new, re-testable version. Run it as a long-running quality gate and regressions get caught before a user sees them.

HIGH

FIX IN SKILL

seen 1x · qlik-data-analyst

Skill invocation fails with no diagnostic information

DRAFTED FIX

1. Verify the skill's definition file has no syntax errors.
2. Add error handling that reports specific failure modes with context.
3. Instrument the invocation layer to log skill-resolution status and argument validation.
4. Re-test with expected arguments.

Occurrences (1): **Run CNYhX** · Failed tool call · "Tool call failed: Skill."

Resolve

Fix with assistant

Qlik AI Workbench
MCP analyzer

Dashboard

MCP Servers

Scans

Compare

Skills

Testing

Collections

Runs

Compatibility

Setup

Environments

Settings

Local / dev mode

Skills > qlik-data-analyst

Skills
3 registered

Search skills...

qlik-ai-readiness-optimizer
GitHub · 2 versions

qlik-data-analyst
GitHub · 8 versions

qlik-sense-app-analysis
GitHub · 2 versions

3 skills · no collisions

Latest (v8)

GitHub

Overview Files Quality Usage Issues 1 Versions Diff

1 open 0 resolved 1 high 0 medium 0 low

Export Markdown Export JSON

Open High Skill Fix in skill seen 1 x · last seen Jul 13, 2026, 7:51 AM

Skill invocation fails with no diagnostic information

The Skill tool call for "qlik-data-analyst" fails with only a generic "Tool call failed: Skill." error message, providing no actionable diagnostic information. This could indicate the skill is not registered, has a loading/syntax error in its definition, encounters an unhandled exception during execution, or receives invalid arguments. The lack of specific error details prevents developers from diagnosing and fixing the root cause.

DRAFT FIX

1. Verify qlik-data-analyst skill exists in the skills registry and its definition file (CLAUDE.md or similar) has no syntax errors. 2. Add comprehensive error handling in the skill to catch and report specific failure modes with context (missing dependencies, invalid arguments, API failures). 3. Instrument the Skill tool invocation layer to log detailed diagnostics: skill resolution status, argument validation results, loading errors, and execution stack traces. 4. Test skill invocation manually with expected arguments to reproduce and isolate the failure point.

Occurrences (1)

Run CNYhX_kn... Jul 13, 2026, 7:51 AM Failed tool call Tool call failed: Skill.

Resolve Fix with assistant

Filed against the skill. A poor run opens a tracked issue on the skill that caused it, drafted fix attached, linked back to the run.

ANSWERS

The questions we get on day one.

Does AI Workbench replace my model provider?



No. It drives your provider's models, or a Qlik Answers assistant, through a real agent loop and measures the result. You bring your own provider key or your Claude subscription for test runs.

What does it cost to run?



The app itself runs locally with no accounts and no cloud. The only spend is your own provider's usage when you execute test runs, and AI Workbench estimates that cost per run so there are no surprises.

How is a run graded?



Automatically, the moment it finishes: answer validation with cited evidence, insight surplus, and an error-forensics inventory. Add expectation grades (Answer, Judge, ROUGE-1, Hygiene, Surplus) when a test defines what good looks like. Every run carries the same grades, which is what makes runs comparable.

Can it really compare Qlik Answers to an MCP agent?



Yes. A Qlik Answers assistant is a first-class run target alongside any model. Run both against the same test and compare turns, tokens, cost, and quality side by side, with a plain-language verdict on the trade-off.

Where does my data go?



Nowhere it does not already go. Servers, credentials, and results stay on your machine. Saved secrets are encrypted at rest and never returned in responses. Test runs call only the providers you configure.

Do I need a Qlik tenant?



Not to analyze MCP servers, run sessions, grade, or compare. A Qlik tenant is only needed if you want to point AI Workbench at a Qlik Cloud MCP server or a Qlik Answers assistant.

Run AI Workbench locally in minutes.

One container, one command. Open it in your browser and connect your first server.

```
docker compose up --build → http://localhost:8080
```

Get started

Read the guide



AI Workbench is the MCP analyzer and session workbench. Measure, test, grade, and compare AI sessions, entirely on your own machine.

© 2026 Qlik. Wherever There's Data.

[Privacy](#) [Terms](#) [Cookies](#)

[Accessibility](#)